

Saharsh Dudhyal

Pune, India | +91 9420781481 | saharshdudhyal@gmail.com | [LinkedIn](#) | [GitHub](#) | [Leetcode](#)

Professional Summary

Computer Science undergraduate specializing in cloud-native software engineering, REST API and microservices development, and AI/ML-driven systems. Experienced building scalable backend services, agentic AI pipelines, and observability platforms using Node.js, Python, Docker, AWS, and modern DevOps practices, with a strong foundation in Agile development, Git-based workflows, CI/CD.

Education

AISSMS Institute of Information Technology, Pune
B.Tech in Computer Science and Engineering — CGPA: 8.13

2023 – 2027

Technical Skills

- **Languages:** C++, Python, JavaScript, SQL
- **Backend & APIs:** Node.js, Express.js, REST APIs, Microservices, WebSockets, API Design
- **Frontend:** React.js, Next.js, HTML5, CSS3
- **AI/ML & GenAI:** LLMs, Retrieval-Augmented Generation (RAG), Prompt Engineering, Agentic AI Workflows, Vector Search
- **Cloud & DevOps:** AWS (EC2, S3), Firebase, Docker, Git, CI/CD
- **Databases & Tools:** SQL, NoSQL (SQLite, Firebase), Splunk, Postman, REST/WebSocket Testing

Experience

Team Lead — Hackshetra 2025 (National-Level Hackathon)

Feb 2025

- Led a 4-member team through a 24-hour national-level hackathon, building **CourseForge**, a custom course-generator.
- Curating disjointed YouTube videos into a structured, pedagogical path to solve the inherent flaws of YouTube learning: **clickbait, content gaps, variable video quality, and lack of active learning.**
- Placed in the **top 50 out of 400+ competing teams.**

Projects

Radium — AI Research Paper Assistant | [Repo](#) | [Live](#)

Jan 2026

Python, RAG, Vector Search, PDFParsing

- Built a Retrieval-Augmented Generation assistant enabling researchers to semantically search and query large academic paper collections instead of manually reading each document.
- Developed a custom PDF parsing and chunking pipeline with vector embeddings, handling document ingestion for 100+ page PDFs.
- Implemented hybrid retrieval combining semantic search with contextual ranking and workspace-based conversational sessions for multi-document Q&A, achieving sub-second retrieval and improved response relevance.

LogPilot — Agentic Observability & Self-Healing Platform | [Repo](#)

May 2026

Node.js, Express.js, Splunk, SQLite, WebSockets, AI/LLMs, MCP

- Built an agentic observability and incident-response platform to solve the lack of unified, real-time monitoring in production systems, reducing manual debugging effort by **70%+**.
- Developed with Node.js, Express.js, Splunk HEC, SQLite, WebSockets, REST APIs, and MCP tool integration to support agentic AI workflows.
- Implemented real-time monitoring across **100% of HTTP requests** and **5+ system metrics** (CPU, memory, event-loop lag, latency, error rate), an AI-powered incident-investigation pipeline with correlation graphs and NLP search, and a self-healing engine with **8+ automated recovery actions** (circuit breaking, rate limiting, garbage collection).
- Packaged as an npm-installable middleware library backed by **143 passing tests**; reduced manual debugging effort by 70%+, cut root-cause analysis time by up to **60%**, and improved incident recovery speed by up to **70%** in simulated production failures.

Certifications

- **AWS Certified Cloud Practitioner (Udemy)** — Foundational cloud concepts, EC2, S3, and cloud-native architecture.

Achievements

- Built and deployed 5+ projects spanning AI/agentic systems, developer tooling, browser extensions, and cloud applications.
- Competed in multiple hackathons, including a Splunk-focused build (LogPilot) applying AI/ML, agentic automation, and cloud-native engineering practices.
- Solved 150+ Data Structures and Algorithms problems on LeetCode (Arrays, Trees, Graphs, DP).
- Developed multiple RAG pipelines enabling semantic search and contextual question answering across large document collections.